



# Artificial Intelligence

## Chapter 4: Beyond Classical Search

Andreas Zell

After the Textbook: Artificial Intelligence,  
A Modern Approach  
by Stuart Russel and Peter Norvig (3<sup>rd</sup> Edition)

## 4. Beyond Classical Search


- Search algorithms considered so far
  - explore the search space systematically
  - keep one or more paths in memory
  - record which alternatives have been explored at each point along the path
- For many problems the path is irrelevant and only the configuration of the goal state is important
- **Local search** algorithms
  - Store only the current node
  - Generally move only to neighbors of this node
  - Do not retain the path to this node

- **Local search** algorithms
  - Are not systematic
  - Require little memory – usually only a constant amount to store a set of current states
  - Often find reasonable solutions in large or even infinite search spaces, where systematic algorithms fail
  - Can be used on **optimization problems**
- **Optimization problems** do not fit the standard model from chapter 3 and often have no goal test or path cost

# 4.1 State-Space Landscape



- Local search algorithms work on state-space landscapes which are defined by the relationship between “location” (state) and “elevation” (value of heuristic cost function or objective function)



- If “elevation” corresponds to cost, the aim is finding the lowest valley – the global minimum
- If “elevation” corresponds to an objective function, the aim is finding the highest peak – the global maximum
- Conversion between both types can be easily done by inserting a minus sign
- A local search algorithm is called **complete** if it always finds a goal if one exists and is called **optimal** if the found goal is a global minimum/maximum

- Basic idea of hill-climbing search (HC):
  - continuously move to neighbor states of increasing value and terminate at a “peak” if no better neighbor than the current state exists
- Only current state with its objective value is stored
- HC does not look beyond immediate neighbors of the current state

**function** HILL-CLIMBING(*problem*) **returns** a state that is a local maximum

*current*  $\leftarrow$  MAKE-NODE(*problem*.INITIAL-STATE)

**loop do**

*neighbor*  $\leftarrow$  a highest-valued successor of *current*

**if** *neighbor*.VALUE  $\leq$  *current*.VALUE **then return** *current*.STATE

*current*  $\leftarrow$  *neighbor*



- Local search algorithms typically use a complete state formulation
- In case of the 8-queen problem each state has 8 queens on the board and stores the row of the queen for each column
- Successor states are all possible states that can be generated by moving a single queen on the board to another row on the same column
- Cost function  $h$  is the number of pairs of queens that are attacking each other
- Global minimum is a state with  $h=0$


## 4.1.1 Hill-Climbing Example 8-Queens



18	12	14	13	13	12	14	14
14	16	13	15	12	14	12	16
14	12	18	13	15	12	14	14
15	14	14	♙	13	16	13	16
♙	14	17	15	♙	14	16	16
17	♙	16	18	15	♙	15	♙
18	14	♙	15	15	14	♙	16
14	14	13	17	12	14	12	18

state (4,3,2,5,4,3,2,3)  $h=17$

with costs for the  
successors obtained by  
moving the queen of this  
column to this row



local minimum with  $h=1$   
each possible successor  
has a higher or equal  
cost




- Hill-climbing is a **greedy local search** since it always moves to the best neighbor state without considering the next moves
- Greedy strategies often make rapid progress at the start but risk to get stuck
- Reasons for getting stuck
  - **Local maxima**: all neighbor states have a worse value than the current state, so no move possible
  - **Plateaux**: neighbor states have same value as current state; no sense of direction; hill climbing might get lost can be alleviated by allowing sideways moves but may end in a loop, therefore requires stopping criterion

## 4.1.1 Greedy Local Search



- Reasons for getting stuck:
  - **Ridges**: Sequence of local maxima that are connected by states with lower value



- **Stochastic hill climbing**: chooses random uphill moves instead of the best; slower convergence but finds better solutions in some landscapes
- **First-choice hill climbing**: instead of generating all moves, subsequently generate moves until a state is found which is better than the current one
  - good strategy for landscapes with extremely high number of neighbor states
- **Random-restart hill climbing**: Repeat hill-climbing runs starting from random initial states until goal is found; trivially complete since an initial state which is also goal may eventually be generated

- **Simulated annealing** (SA) works similar to hill-climbing but also allows “downhill” moves
- Mimicks the annealing process in metallurgy, which starts at high temperature  $T$  and allows to reach a low energy state by gradual cooling
- A random move is generated at each iteration
- If it improves the objective function the move is always accepted
- If the objective function is worsened the move is accepted with probability  $p$
- $p$  is exponentially reduced with the “badness” of the move and also decreases with the temperature  $T$

## 4.1.2 Simulated Annealing

- Simulated annealing starts with high  $T$  which is reduced over time; bad moves are more likely to be allowed at the start and become more unlikely
- If  $T$  is lowered slowly enough SA finds the global optimum with probability 1

**function** SIMULATED-ANNEALING(*problem*, *schedule*) **returns** a solution state

**inputs:** *problem*, a problem

*schedule*, a mapping from time to “temperature”

*current*  $\leftarrow$  MAKE-NODE(*problem*.INITIAL-STATE)

**for**  $t = 1$  **to**  $\infty$  **do**

$T \leftarrow$  *schedule*( $t$ )

**if**  $T = 0$  **then return** *current*

*next*  $\leftarrow$  a randomly selected successor of *current*

$\Delta E \leftarrow$  *next*.VALUE  $-$  *current*.VALUE

**if**  $\Delta E > 0$  **then** *current*  $\leftarrow$  *next*

**else** *current*  $\leftarrow$  *next* only with probability  $e^{\Delta E/T}$

- Instead of only one state **local beam search** keeps track of  $k$  states in parallel
- It is initialized with  $k$  random states and at each step generates all possible successors of the  $k$  current states
- If one of the successors is the goal, local beam search stops; otherwise the  $k$  best successors are selected as the new current states
- Can suffer from lack of diversity, since the  $k$  states quickly become concentrated at a small region turning local beam search into a form of hill-climbing
- **Stochastic beam search** alleviates this by randomly selecting the  $k$  successors with a probability corresponding to their objective function values

- **Genetic algorithms** are a variant of stochastic beam search based on principles of natural evolution
- States are represented by **individuals** within a **population**
- Each iteration new individuals are generated by **crossover** of two randomly selected parents within the current population
- The probability to be selected as parent corresponds to the objective function value respectively **fitness function** value of the individual
- With a low probability the individuals are slightly changed by **mutation** after crossover
- Lecture in summer semester “**Evolutionäre Algorithmen**“

## 4.1.4 Genetic Algorithms

**function** GENETIC-ALGORITHM(*population*, FITNESS-FN) **returns** an individual

**inputs:** *population*, a set of individuals

FITNESS-FN, a function that measures the fitness of an individual

**repeat**

*new\_population*  $\leftarrow$  empty set

**for**  $i = 1$  **to** SIZE(*population*) **do**

*x*  $\leftarrow$  RANDOM-SELECTION(*population*, FITNESS-FN)

*y*  $\leftarrow$  RANDOM-SELECTION(*population*, FITNESS-FN)

*child*  $\leftarrow$  REPRODUCE(*x*, *y*)

**if** (small random probability) **then** *child*  $\leftarrow$  MUTATE(*child*)

add *child* to *new\_population*

*population*  $\leftarrow$  *new\_population*

**until** some individual is fit enough, or enough time has elapsed

**return** the best individual in *population*, according to FITNESS-FN

**function** REPRODUCE(*x*, *y*) **returns** an individual


**inputs:** *x*, *y*, parent individuals

$n \leftarrow$  LENGTH(*x*);  $c \leftarrow$  random number from 1 to  $n$

**return** APPEND(SUBSTRING(*x*, 1,  $c$ ), SUBSTRING(*y*,  $c + 1$ ,  $n$ ))



## 4.1.4 Genetic Algorithms Example





- Almost all algorithms (with exception of hill-climbing and simulated annealing) discussed so far can not handle continuous search spaces
- Since the development of calculus by Newton and Leibniz in the 17<sup>th</sup> century numerous algorithms for continuous optimization have been published
- As example consider the problem of placing three airports in Romania, so that the sum of the squared distances of these airports to their closest cities is minimal
- The state space is then defined by the coordinates of the airports:  $(x_1, y_1), (x_2, y_2), (x_3, y_3)$



- The resulting search space contains 6 **variables**
- Moving around in the search space corresponds to moving one or more airports
- Let  $C_i$  be the set of closest cities for airport  $i$ , then the objective function for the neighborhood of the current state, where  $C_i$  remains constant, is:

$$f(x_1, y_1, x_2, y_2, x_3, y_3) = \sum_{i=1}^3 \sum_{c \in C_i} \left( (x_i - x_c)^2 + (y_i - y_c)^2 \right)$$

- This expression is correct locally but not globally because the relationship of a city to the sets  $C_i$  is a discontinuous function of the state



- A way to avoid continuous problems is to **discretize** the neighborhood of each state
- In the example this can be done by moving only one airport by a fixed amount  $\pm\delta$  giving 12 possible successors for each state
- The resulting problem can be solved by any previously described local search algorithm
- Simulated annealing and stochastic hill climbing can also be applied with randomly generated vectors of length  $\delta$

## 4.2 Using gradient information



- Many methods attempt to navigate in the landscape using the **gradient** which is the magnitude and direction of the steepest slope
- In our example the gradient is:

$$\nabla f = \left( \frac{\partial f}{\partial x_1}, \frac{\partial f}{\partial y_1}, \frac{\partial f}{\partial x_2}, \frac{\partial f}{\partial y_2}, \frac{\partial f}{\partial x_3}, \frac{\partial f}{\partial y_3} \right)$$

- Some problems can be solved by solving the equation  $\nabla f = 0$
- In many cases this equation can not be solved in closed form, as in our example, since the gradient depends on  $C_i$  which is discontinuous

## 4.2 Using gradient information

- We can still compute the gradient locally for example for airport 1 and constant  $C_1$

$$\frac{\partial f}{\partial x_1} = 2 \sum_{c \in C_1} (x_i - x_c)$$

- Given a locally correct gradient we can perform **steepest ascent hill-climbing** by updating the state according to

$$\mathbf{x} \leftarrow \mathbf{x} + \alpha \nabla f(\mathbf{x})$$

- $\alpha$  is the **step size**, which is a small constant
- If  $f$  is not differentiable, an **empirical gradient** can be obtained by evaluating the response to small changes in each variable

- Determining optimal  $\alpha$  is complicated. Too small  $\alpha$  leads to many too small steps. Too large  $\alpha$  leads to overshooting the optimum
- **Line search** tries to overcome this by extending the current gradient direction until  $f$  starts to worsen again
- For many problems the most efficient method is the **Newton-Raphson** algorithm, a general technique for finding roots of functions, solving the equation  $g(x)=0$  by iteratively computing new estimates for the root  $x$  by:

$$x \leftarrow x - g(x) / g'(x)$$

## 4.2 Newton-Raphson algorithm

- For optimizing  $f$  we need to find an  $\mathbf{x}$  so that  $\nabla f = 0$
- By setting  $\nabla f = g(\mathbf{x})$  we obtain:  

$$\mathbf{x} \leftarrow \mathbf{x} - H_f^{-1}(\mathbf{x}) \cdot \nabla f(\mathbf{x})$$
- $H_f(\mathbf{x})$  is the Hessian matrix of second derivatives with the elements  $\partial^2 f / \partial x_i \partial x_j$
- For the airport example  $H_f(\mathbf{x})$  is zero for all off-diagonal elements and the diagonal elements for airport  $i$  are twice the number of cities in  $C_i$
- An update step then moves the airport  $i$  directly to the centroid of  $C_i$



## 4.2 Newton-Raphson algorithm

- For high-dimensional problems computing the  $n^2$  elements of  $H_f(\mathbf{x})$  is expensive, therefore many approximate versions of the method have been developed
- As local search algorithm gradient based algorithms also suffer from local optima, ridges, plateaux
- Random restart or simulated annealing can still be used

- A final topic to this is **constrained optimization**, which are problems whose solutions must satisfy some hard constraints on the values of the variables
- The airports in the example might be constrained to be inside Romania and on dry land
- The difficulty of such problems depends on the nature of the constraints
- **Linear programming** is the best known category, in which constraints must be linear inequalities forming a **convex set** and the objective must also be linear. Problems of this type can be solved in polynomial time on the number of variables




- So far we assumed that the environment is fully observable and that the agent knows the effect of each action
- When the environment is either partially observable or nondeterministic, **percepts** become useful
- In partially observable environments percepts help to narrow down the set of possible states the agent might be in
- In nondeterministic environments, percepts tell the agent which possible outcome of the action has actually occurred

- Percepts can not be determined in advance but the agents actions will depend on future percepts
- Thus the solution to a problem is not a sequence of actions but a **contingency plan (strategy)**
- As example we will use the vacuum cleaner world with the actions *Left, Right, Suck*
- In the **erratic vacuum world** *Suck* is redefined as:
  - When applied to a dirty square, the square is cleaned and an adjacent square sometimes is cleaned up, too.
  - When applied to a clean square, sometimes dirt is deposited on that square

## 4.3 Erratic Vacuum World



- The eight possible states of the erratic vacuum world – states 7 and 8 are goal states



## 4.3 Erratic Vacuum World


- To provide a precise formulation we need to generalize the **transition model** of Chapter 3
- Instead of defining the transition model by a RESULT function which returns a single state we use a RESULT function which returns a set of possible outcomes
- For example *Suck* in state 1 returns the set {5,7}
- We also need to generalize the concept of solution, since, for example, if we start in state 1 there is no single sequence of actions to solve the problem; instead we need a contingency plan like:  
     *[Suck, if State=5 then [Right,Suck] else []]*

- Thus solutions for nondeterministic problems can contain nested **if-then-else** statements
- This results in trees rather than sequences allowing to select actions based on contingencies arising during execution
- Many problems in the real, physical world are such contingency problems, since exact prediction is impossible
- In a deterministic environment branches in a tree search are introduced by the agents' own choices in each state. We call such nodes **OR nodes**

## 4.3 AND-OR Search Trees



- In a **nondeterministic environment** branching is also introduced by the environment's choice of the outcome of each action. We call these **AND nodes**
- The two kind of nodes alternate leading to an **AND-OR tree**
- The bold path corresponds to the plan given before





- A solution for an AND-OR search problem is a subtree that:
  1. has a goal node at every leaf
  2. specifies one action at each of its OR nodes
  3. includes every outcome branch at each of its AND nodes
- The resulting plan uses if-then-else notation to handle AND branches
- If there are more than two branches at a node a **case** construct might be better

- Modifying the basic problem-solving agent from 3.1 to execute such solutions is straightforward
- One might also consider a somewhat different agent design which acts before it has found a guaranteed plan and deals with contingencies as they arise
- This type of interleaving search and execution is also useful for exploration problems
- One key aspect is the way to deal with cycles which often arise in nondeterministic problems
- The following algorithm returns failure if the current state is identical to a state on the path from the root

## 4.3 AND-OR Search Trees

- This does not mean there is no solution, only that if there is a non-cyclic solution it must be reachable from an earlier state

---

**function** AND-OR-GRAPH-SEARCH(*problem*) **returns** *a conditional plan, or failure*  
OR-SEARCH(*problem*.INITIAL-STATE, *problem*, [])

---

**function** OR-SEARCH(*state*, *problem*, *path*) **returns** *a conditional plan, or failure*  
**if** *problem*.GOAL-TEST(*state*) **then return** the empty plan  
**if** *state* is on *path* **then return** failure  
**for each** *action* **in** *problem*.ACTIONS(*state*) **do**  
    *plan*  $\leftarrow$  AND-SEARCH(RESULTS(*state*, *action*), *problem*, [*state* | *path*])  
    **if** *plan*  $\neq$  failure **then return** [*action* | *plan*]  
**return** failure

---

**function** AND-SEARCH(*states*, *problem*, *path*) **returns** *a conditional plan, or failure*  
**for each**  $s_i$  **in** *states* **do**  
    *plan*<sub>*i*</sub>  $\leftarrow$  OR-SEARCH( $s_i$ , *problem*, *path*)  
    **if** *plan*<sub>*i*</sub> = failure **then return** failure  
**return** [**if**  $s_1$  **then** *plan*<sub>1</sub> **else if**  $s_2$  **then** *plan*<sub>2</sub> **else** ... **if**  $s_{n-1}$  **then** *plan* <sub>$n-1$</sub>  **else** *plan* <sub>$n$</sub> ]


---

- AND-OR graphs can also be explored by breadth-first or best-first methods
- For this the concept of heuristic functions must be modified to estimate the cost of a contingent solution rather than a sequence
- The notion of admissibility carries over and there is an analog version of the  $A^*$  algorithm
- Now consider the **slippery vacuum world** which is identical to the erratic vacuum world except that movement actions may fail, e.g. *Right* in the state 1 thus leads to the state set  $\{1,2\}$

## 4.3 AND-OR Search Trees



- There are no longer any acyclic solutions from state 1
- AND-OR-graph-search would return failure
- However there is a cyclic solution which is to keep trying *Right* until it works




- We can express this solution by adding a label denoting some portion of the plan and use it later  
[*Suck, L1 : Right* **if** *State=5* **then** *L1* **else** *Suck*]
- A cyclic plan may be considered a solution provided that every leaf is a goal state and that a leaf is reachable from every point in the plan
- Given the definition of a cyclic solution an agent executing this solution will only eventually reach the goal with probability 1, if the nondeterminism is caused by a stochastic process, and not by a hidden property of the environment which prevents some state transitions

- In problems with partial observability, the agent's percepts are insufficient to pin down the exact state
- An action may in this case lead to one of several states even if the environment is deterministic
- The key concept for such problems is the **belief state** which represents the agent's current belief about possible physical states it might be in, given the sequence of actions and percepts
- First we consider a **sensorless** agent, whose percepts provide no information at all
- Such problems with sensorless agents, are also called **conformant** problems

## 4.4.1 Searching with no Observation


- Sensorless Agents can still often act successful and actually be advantageous since they don't rely on costly and potentially unreliable sensor information
- As example we assume a sensorless vacuum world where the agent knows the geography of its world but not its position or the distribution of dirt
- The initial state then can be any element of the set  $\{1,2,3,4,5,6,7,8\}$





## 4.4.1 Searching with no Observation


- If we move to the *Right* we reduce the set of possible states to  $\{2,4,6,8\}$



- Using *Suck* we can reduce the set further to  $\{4,8\}$



- Using the sequence  $[Right, Suck, Left, Suck]$  guarantees to reach in the goal state 7 regardless of the initial state



- For this problem we say the agent can **coerce** the world into state 7.
- To solve sensorless problems we search in the space of belief states rather than in the space of physical states.
- The belief space of the problem is fully observable because the agent always knows its belief state.
- The solution, if it exists, is always a sequence of actions, since the percepts are always empty and therefore completely predictable, so there are no contingencies to plan for.



- For the physical problem  $P$  we can now define the corresponding sensorless problem as follows:
  - **Belief states**: the entire belief state space containing every possible set of physical states  $P$  (up to  $2^N$  belief states if  $P$  has  $N$  possible states)
  - **Initial state**: Typically the set of all states in  $P$
  - **Actions**: The actions in a belief state are the union of the actions of the physical states

$$\text{ACTIONS}(b) = \bigcup_{s \in b} \text{ACTIONS}_p(s)$$

(we here ignore forbidden actions)

- Transition model:

- The result of an action is the set of all physical states resulting from performing the action on all physical states in the current belief state

- For deterministic actions the set is:

$$b' = \text{RESULT}(b, a) = \{s' : s' = \text{RESULT}_p(s, a) \text{ and } s \in b\}$$

- For nondeterministic actions the set is:

$$b' = \text{RESULT}(b, a) = \{s' : s' \in \text{RESULT}_p(s, a) \text{ and } s \in b\}$$


$$= \bigcup_{s \in b} \text{RESULT}_p(s, a)$$

- The process of generating a new belief state after action is called the prediction step  $\text{PREDICT}_p(b, a)$


## 4.4.1 Searching with no Observation



- For deterministic actions  $b'$  is never larger than  $b$ , but for nondeterministic actions  $b'$  might be larger
- (a) Predicting the next belief state for the sensorless vacuum world with the deterministic action *RIGHT*
- (b) Predicting the same belief state for the same action in the slippery vacuum world



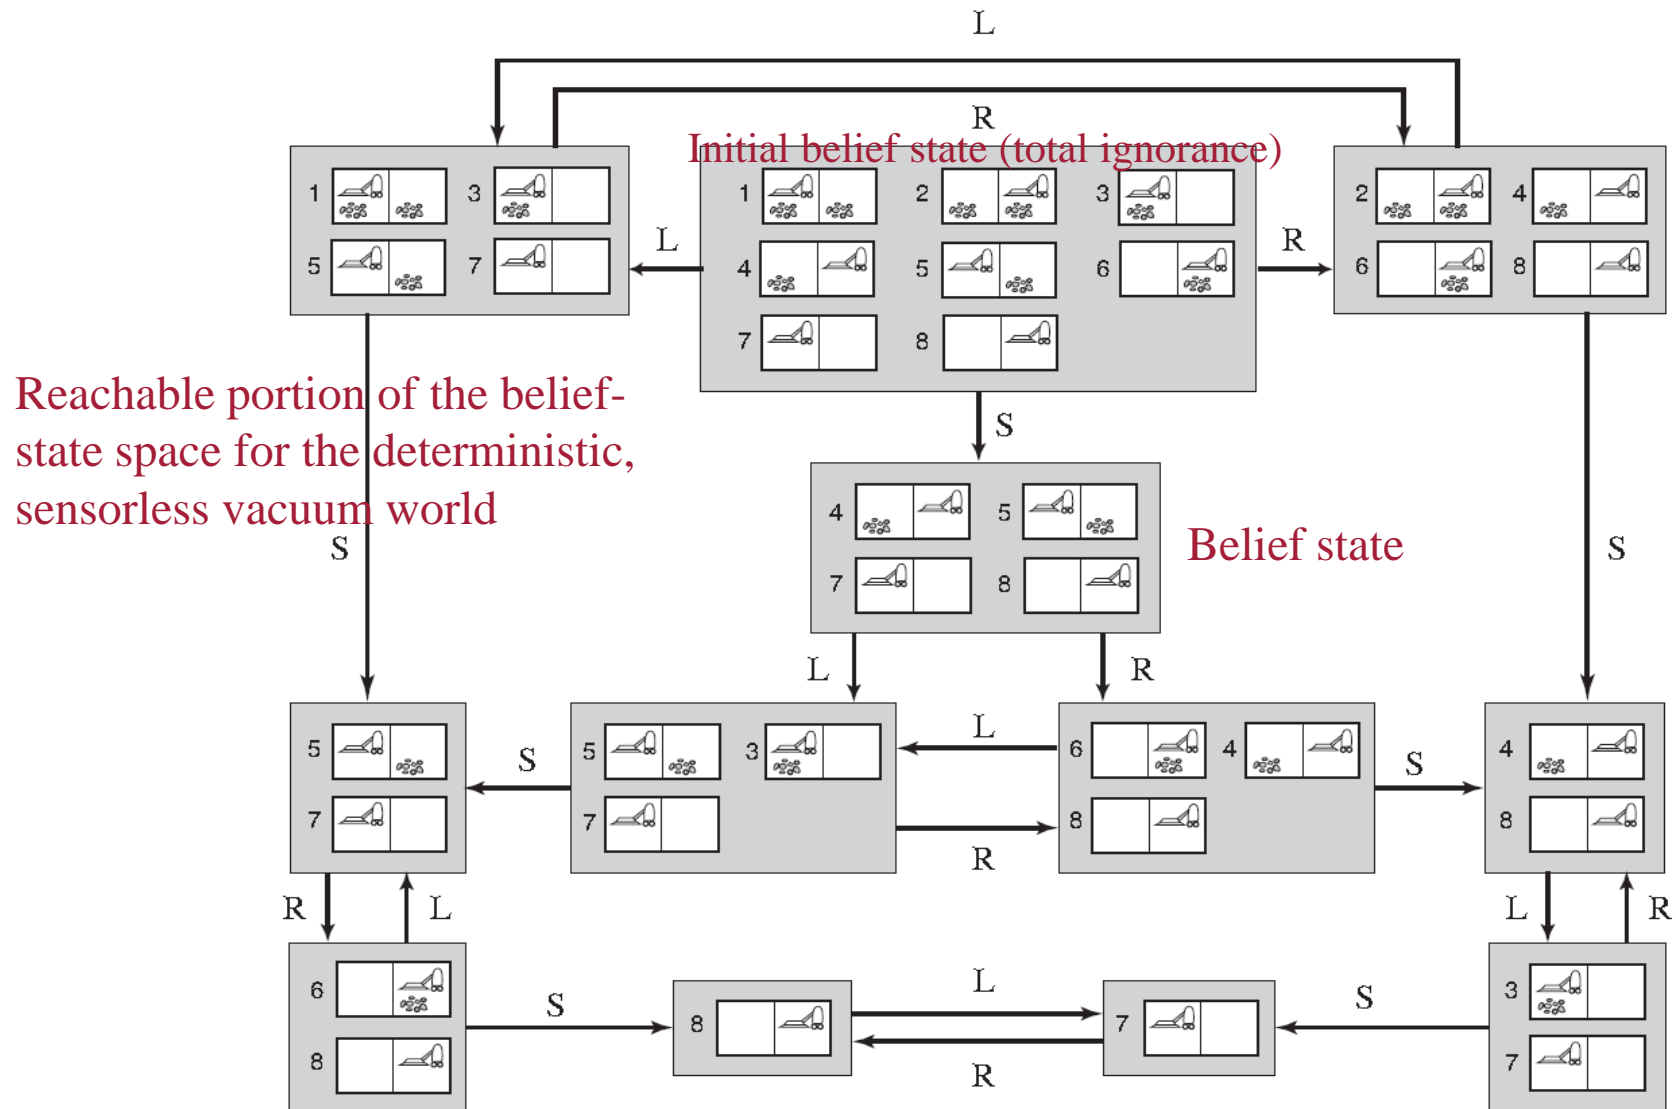
(a)



(b)

- **Goal test:**
  - Goal states in the belief space are only states where all physical states satisfy the  $\text{GOAL-TEST}_P$
  - The agent may accidentally arrive at the goal earlier but it will not know this
- **Path cost:**
  - We assume that any action has identical cost for all physical states in the same belief state
- The definitions allow an automatic construction of the belief state problem
- Any previously described search algorithm can be applied

## 4.4.1 Searching with no observation



## 4.4.1 Searching with no Observation



- Even with pruning, sensorless problem solving is seldom feasible using the algorithms described so far because of the size of the belief space
- For example the initial belief state of a 10x10 vacuum world contains  $100 \times 2^{100}$  or  $10^{32}$  physical states
- A solution to this is to represent the belief state in a more compact description using formal representation schemes (Chapter 7)
- Another solution is to avoid treating belief states as black boxes
- Instead we can use incremental belief state search which looks inside belief states




- For general partially observable problems we have to specify how the environment generates percepts for the agent
- Using a `PERCEPT(s)` function we can extend the method to general partial observable problems
- `PERCEPT(s)` returns the percept of the agent for a given state  $s$
- `PERCEPT(s)` returns a set of possible percepts for cases with nondeterministic sensing
- Fully observable problems have  $\text{PERCEPT}(s) = s$
- Sensorless problems have  $\text{PERCEPT}(s) = \text{null}$

## 4.4.2 Searching with Observations



- For partially observable problems the same percept could have been produced by several states, e.g. the percept  $[A, \textit{Dirty}]$  can be produced by the states  $\{1, 3\}$  in the vacuum world.



- We can keep ACTION, STEP-COST, and GOAL-TEST from the underlying problem just as for sensorless problems
- We have to incorporate the percepts into the transition model for the belief states

- We divide the transition from one belief state to an other into three stages: prediction, observation-prediction, update
- **Prediction**: is the same as for sensorless problems; given an action  $a$  in belief state  $b$  the predicted belief state is  $b' = \text{PREDICT}(b, a)$
- **Observation Prediction**: determines the set of percepts  $o$  that could be observed in  $b'$ :

$$\text{POSSIBLE-PERCEPTS}(b') = \{o : o = \text{PERCEPT}(s) \text{ and } s \in b'\}$$



- **Update:** determines for each possible percept the belief state that would result from the percept. The new belief state  $b_o$  is the set of states in  $b'$  that could have produced the percept


$$b_o = \text{UPDATE}(b', o) = \{s : o = \text{PERCEPT}(s) \text{ and } s \in b'\}$$

- The updated belief state  $b_o$  can not be larger than the predicted belief  $b'$  state since observations can only help to reduce the uncertainty compared to the sensorless case
- For deterministic sensing, the belief states for the different percepts will be disjoint forming a partition of the original predicted belief state

## 4.4.2 Searching with Observations



- Transition in the **deterministic case**: *Right* is applied in the initial belief state resulting in a new belief state with two possible physical states depending on the percepts:  $[B, Dirty]$ ,  $[B, Clean]$
- Transition in the **slippery world**: Applying *Right* gives a new belief state with four physical states leading to three belief states depending on the percepts:  $[B, Dirty]$ ,  $[A, Dirty]$ ,  $[B, Clean]$



- Putting the three stages together, we obtain the possible belief states resulting from a given action and subsequent possible percept by:

$$\text{RESULTS}(b, a) = \left\{ b_o : \begin{array}{l} b_o = \text{UPDATE}(\text{PREDICT}(b, a), o) \text{ and} \\ o \in \text{POSSIBLE-PERCEPTS}(\text{PREDICT}(b, a)) \end{array} \right.$$

- Given this formulation AND-OR search can be applied directly to find a solution
- As in the sensorless case AND-OR search treats belief states as black boxes.
- The tree can be pruned by checking for subsets or supersets of the current state in the already generated belief states.

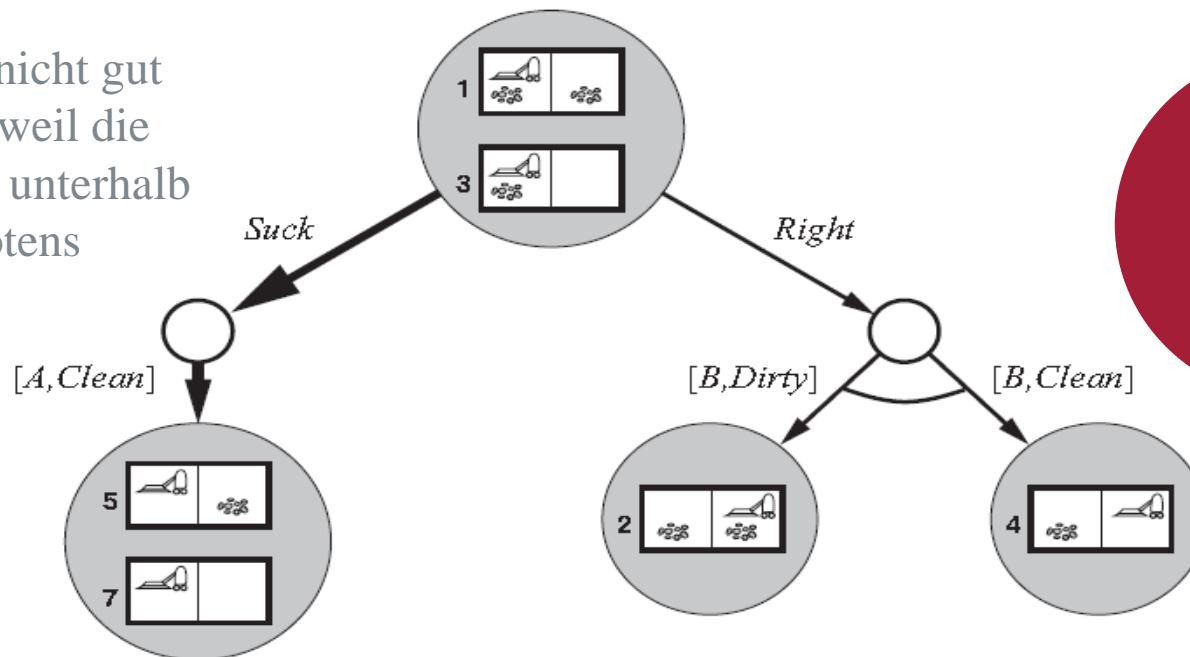
## 4.4.3 Solving partially observable problems



- First level of an AND-OR search tree for the initial percept  $[A, \text{Dirty}]$
- The solution is the conditional plan

$[Suck, Right \text{ if } BState = \{6\} \text{ then } Suck \text{ else } []]$

(Lösung passt nicht gut zum Graphen, weil die if-Abfrage erst unterhalb des linken Knotens auftaucht.)





- An agent for partially observable problems is similar to the simple problem solving agent using the following elements
- The main differences are:
  - The solution is a conditional plan instead of a sequence
  - The agent has to maintain its belief state as it performs actions and receives percepts
- Maintaining the belief state is similar to the update step in the search phase and can be given a belief state  $b$ , an action  $a$ , and an observation  $o$  written as:

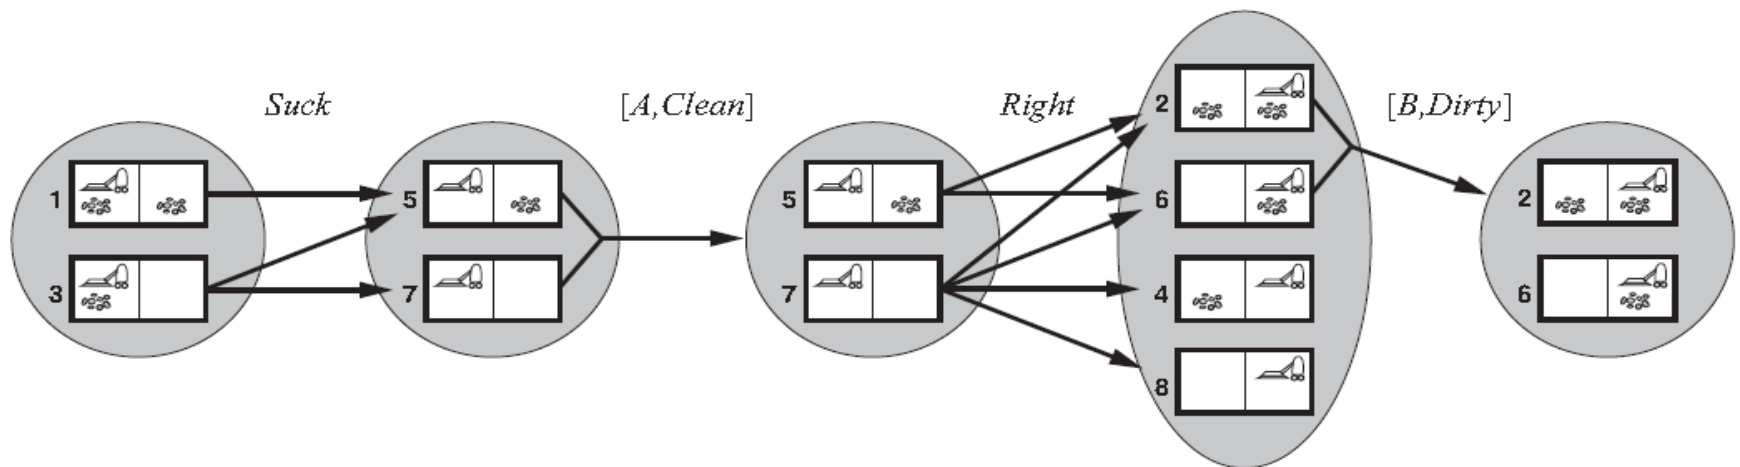
$$b' = \text{UPDATE}(\text{PREDICT}(b, a), o)$$



## 4.4.4 Partially observable environments



- Consider a slightly modified vacuum world where a square can become dirty again unless the agent is cleaning it actively at the time (**kindergarten** world).
- The resulting belief states for two update-prediction cycles for an agent in the **kindergarten vacuum world** are






- Maintaining the belief state is a core function for any intelligent system in a partially observable environment – which includes the vast majority of real world problems
- This function is also called: **monitoring, filtering, or state estimation**
- The update-prediction step has to be as fast as the percepts are coming in, otherwise the agent falls behind
- In more complex environments an exact computation becomes unfeasible and approximations have to be used


## 4.4.4 Partially observable environments



- **Localization** is a problem with partial observations
- E.g. consider determining the possible positions of an erratic robot in a labyrinth with 4 sensors indicating the state of the cells in its 4-neighborhood




(a) Possible locations of robot after  $E_1 = \text{NSW}$




Move  
(nondeterministic)

Robot must be  
here, after  
 $E_1$ , Move,  $E_2$ .



(b) Possible locations of robot After  $E_1 = \text{NSW}, E_2 = \text{NS}$





- So far we have examined **offline search** algorithms which compute a complete solution before acting
- In contrast we now consider **online search** agents which **interleave** computation and action
- Online search is a good idea in dynamic and semidynamic domains
- It is also helpful in nondeterministic domains because it allows the computational effort on the contingencies that actually arise rather than those that might happen
- Online search is necessary for unknown environments where the agent does not know what states exist or what an action does (**exploration problems**)

- A typical example for exploration problems is a robot that is placed in a new building where it has to build a map that it can use to get from  $A$  to  $B$
- First we assume a deterministic and fully observable environment with
  - $ACTIONS(s)$ : returns a list of all allowed actions in state  $s$
  - Step-cost function  $c(s, a, s')$
  - $GOAL-TEST(s)$
- NOTE: We can not determine  $RESULT(s, a)$  except by actually doing action  $a$  in state  $s$

- We may also have an admissible heuristic  $h(s)$  estimating the cost to goal  $G$
- The objective is typically to reach a goal state with minimal cost
- A common performance measure is the **competitive ratio** which compares the actual cost with the best cost that would be achievable if the agent knows the whole state space in advance
- The best achievable cost ratio can be infinite if some actions are **irreversible** and the search reaches a **dead-end state** from which no goal is reachable



- No algorithm can avoid dead-ends in all state spaces
- For an agent both dead-end spaces look identical if it has visited only S and A so it has to take the same decision for both spaces: it will fail in one of them



- We now assume a **safely explorable** state space
- This means that a goal state is reachable from any reachable state
- State spaces with reversible actions like 8-queens or mazes can be viewed as undirected graphs and are safely explorable
- Even for safely explorable environments no bounded competitive ratio can be guaranteed if there are paths of unbounded cost



- After each action the agent receives a percept telling it the reached state
- This can be used to augment the agent's map of the environment
- The map can be used to decide the next action
- This interleaving is different from offline agents which can immediately expand the next node
- An online agent can only expand the node it physically occupies
- To avoid traveling all over the search tree to expand the next node it is better to expand the nodes in local order

- Depth-first-search works locally since the next node is always a child of the current node

**function** ONLINE-DFS-AGENT( $s'$ ) **returns** an action

**inputs:**  $s'$ , a percept that identifies the current state

**persistent:** *result*, a table indexed by state and action, initially empty

*untried*, a table that lists, for each state, the actions not yet tried

*unbacktracked*, a table that lists, for each state, the backtracks not yet tried

$s$ ,  $a$ , the previous state and action, initially null

**if** GOAL-TEST( $s'$ ) **then return** *stop*

**if**  $s'$  is a new state (not in *untried*) **then** *untried*[ $s'$ ]  $\leftarrow$  ACTIONS( $s'$ )

**if**  $s$  is not null **then**

*result*[ $s$ ,  $a$ ]  $\leftarrow s'$

add  $s$  to the front of *unbacktracked*[ $s'$ ]

**if** *untried*[ $s'$ ] is empty **then**

**if** *unbacktracked*[ $s'$ ] is empty **then return** *stop*

**else**  $a \leftarrow$  an action  $b$  such that *result*[ $s'$ ,  $b$ ] = POP(*unbacktracked*[ $s'$ ])

**else**  $a \leftarrow$  POP(*untried*[ $s'$ ])


$s \leftarrow s'$

- The agent stores its map in a table  $\text{RESULT}[s,a]$  that records the state resulting from executing action  $a$  in state  $s$
- When the agent reaches an action that has not been tried in the state it executes this action
- When all actions of the current state have been tried the agent has to physically backtrack instead of just dropping the node
- **Backtracking** means going back to the state from which the agent most recently entered the current state

## 4.5 Online DFS Agent



- When put in a maze the agent will in the worst case traverse every link twice
- For exploration this is optimal
- For finding a goal the competitive ratio can be arbitrarily bad if the goal is right next to the initial state
- An online variant of iterative deepening solves this problem
- **ONLINE-DFS-AGENT** only works in state spaces with reversible actions




- Hill-climbing search is already an online search algorithm because:
  - It only keeps the current state in memory, which can be identical to the physical state of the agent
  - It only searches locally, therefore it only visits reachable neighbor states
- In its simplest form it is not very useful since it may leave the agent at a local maximum
- Random restarts can not be used, because the agent can not teleport to a random new state

## 4.5 Random Walk



- Instead of random restarts **random walk** can be used to explore the environment
- It can be proven that random walk will eventually find a goal if the state space is finite
- The search can be extremely inefficient
- Consider a state space where backward steps are twice as likely as forward progress
- Random walk will take exponentially many steps to find the goal




- Augmenting random walk with memory can be very effective
- The basic idea is to store a current best estimate of the cost to reach the goal from each state  $s$  that has been visited:  $H(s)$
- Without information  $H(s)$  starts as the heuristic  $h(s)$
- $H(s)$  is updated as the agent gains experience in the state space

## 4.5 Learning Real-Time A\*



- Consider a simple example in a 1D state space




- The estimated cost  $H(s)$  is denoted in each state
- The agent seems to be stuck in a local minimum (shaded area)
- The estimated cost to reach the goal through a neighbor  $s'$  is the cost  $c(s, a, s')$  to get to  $s'$  plus the estimated cost  $H(s')$  to get to the goal from  $s'$


$$\text{cost\_estimate}(s, a, G) = c(s, a, s') + H(s')$$



## 4.5 Learning Real-Time A\*




- For this state, there are two possible actions
  - Left with the estimated cost  $1+9$
  - Right with the estimated cost  $1+2$
- The best move seems to move to the right
- $H(s)=2$  was overly optimistic and has to be updated to  $H(s)=3$






## 4.5 Learning Real-Time A\*




copied from previous page



- Continuing this, the agent moves back and forth twice more updating  $H$  each time



## 4.5 Learning Real-Time A\*



- The agent can now leave the local minimum to the right
- An agent implementing this scheme is called **learning real-time A\* (LRTA\*)**
- Like ONLINE-DFS-AGENT it builds a map of the environment in the result table
- It updates the cost estimate of the state that it has just left and then chooses the apparently best move as next move

- Actions that have not been tried are always assumed to lead to the goal with minimal cost  $h(s)$
- This **optimism under uncertainty** encourages the agent to explore new possible actions
- An LRTA\* agent is guaranteed to find a goal in any finite safely explorable environment
- Unlike A\* it is not complete for infinite state spaces: there are cases where it can be led infinitely astray
- It can explore an environment with  $n$  states in  $O(n^2)$

## 4.5 Learning Real-Time A\*

**function** LRTA\*-AGENT( $s'$ ) **returns** an action

**inputs:**  $s'$ , a percept that identifies the current state

**persistent:** *result*, a table, indexed by state and action, initially empty

*H*, a table of cost estimates indexed by state, initially empty

$s$ ,  $a$ , the previous state and action, initially null

**if** GOAL-TEST( $s'$ ) **then return** *stop*

**if**  $s'$  is a new state (not in  $H$ ) **then**  $H[s'] \leftarrow h(s')$

**if**  $s$  is not null

$result[s, a] \leftarrow s'$

$H[s] \leftarrow \min_{b \in \text{ACTIONS}(s)} \text{LRTA}^*\text{-COST}(s, b, result[s, b], H)$

$a \leftarrow$  an action  $b$  in  $\text{ACTIONS}(s')$  that minimizes  $\text{LRTA}^*\text{-COST}(s', b, result[s', b], H)$

$s \leftarrow s'$

**return**  $a$

**function** LRTA\*-COST( $s, a, s', H$ ) **returns** a cost estimate

**if**  $s'$  is undefined **then return**  $h(s)$

**else return**  $c(s, a, s') + H[s']$

- The initial ignorance of online search provides several opportunities for learning
  - The agent can learn a “map” of the environment - the outcome of each action in each state
  - The agent can acquire more accurate estimates of the cost for each state by using local update rules
- Once exact values are known, optimal decisions can be taken by simply moving to the lowest cost successor; pure hill climbing is then the optimal strategy



- Local search methods like **hill climbing** operate on complete state formulations, keeping only a small number of nodes in memory
- Several stochastic algorithms including **simulated annealing** have been developed
- Many local search methods also apply to continuous spaces
- **Linear programming** and **convex optimization** problems obey certain restrictions and allow polynomial-time algorithms

- A **genetic algorithm** is a **stochastic hill climbing** search which holds a **population** of states and generates new states by **mutation** and **crossover**
- **AND-OR search** can generate **contingency plans** to reach the goal in **nondeterministic** environments
- **Belief states** represent the set of possible physical states in partially observable environments
- Standard search algorithms can be directly applied to belief states to solve **sensorless** problems
- Belief AND-OR search can solve partially observable problems



- **Incremental** algorithms that construct solutions state by state within a belief state are often more efficient
- In **exploration problems** the agent has no knowledge about the states and actions
- **Online search agents** can build a map and find a goal if one exists
- Updating heuristic estimates provides an efficient method to escape from local minima