



# Artificial Intelligence

## Chapter 8: First-Order Logic

Andreas Zell

After the Textbook: Artificial Intelligence,  
A Modern Approach  
by Stuart Russel and Peter Norvig (3<sup>rd</sup> Edition)

- One drawback of propositional logic is that it can't describe environments with many objects efficiently
- For example, in the Wumpus world, for each square there needs to be a rule for pits, such as
  - $B_{1,1} \Leftrightarrow (P_{1,2} \vee P_{2,1})$
- In **natural languages**, it is easy to say “Squares adjacent to pits are breezy”, even for a grid of **infinite** size



- Natural languages like English or German:
  - are far more **expressive** than propositional logic
  - serve as medium for **communication**
  - depend heavily on the **context**
  - can be **ambiguous**
- We can adapt propositional logic (PL) by borrowing advantages from natural languages, while keeping the **declarative**, **compositional** semantics that is **context-independent** and **unambiguous**.

- Obvious elements of natural language are:
  - nouns and noun phrases that refer to **objects**
  - verbs and verb phrases that refer to **relations**
  - some relations that are **functions**
- Examples are:
  - Objects: people, houses, colors, time, ...
  - Relations:
    - unary relations or **properties**: red, round, prime
    - n-ary relations: brother of, bigger than, part of, owns, ...
  - Functions: father of, best friend, one more than, ...

- First-order logic (FOL) is built around **objects** and **relations**
- It can also express facts about **some** or **all** objects in the universe
- Difference in **ontological commitment**
  - PL assumes that there are **facts that can either be true or false**, and each model assigns true or false to each fact via its propositional symbol
  - FOL also assumes that there are **objects with certain relations** among them that do or do not hold. Its formal models are more complicated.

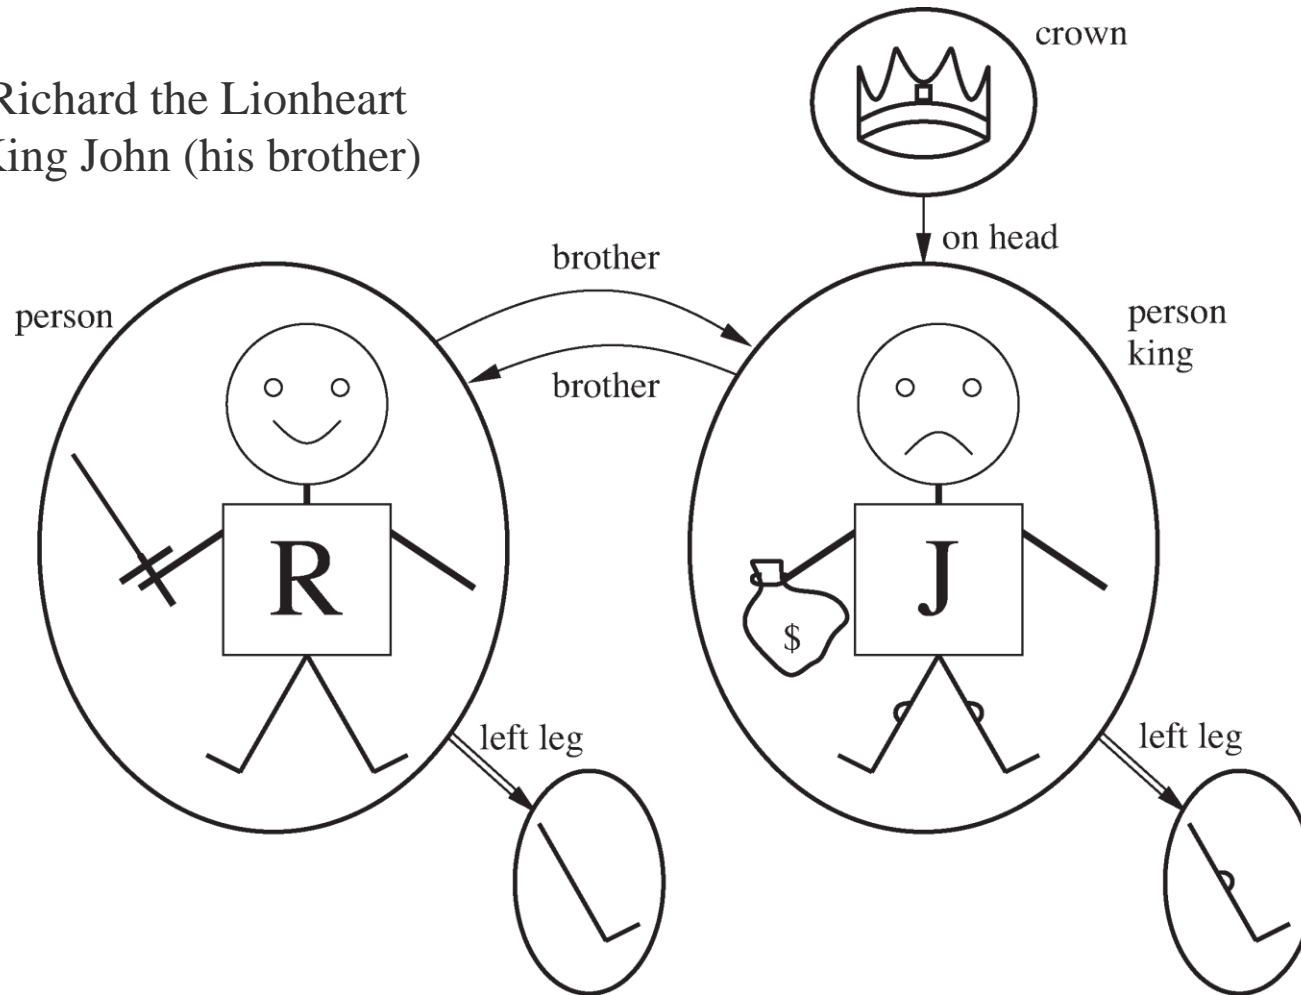


- The domain of a model in FOL is the set of objects or **domain elements** that it contains
- It must be **nonempty**, every possible world must contain at least one object
- It doesn't matter what these objects are
- In the following examples, we use **named objects**, i.e., we refer to objects using their names
- We could also use pictures for it

## 8.2 Syntax and Semantics of First-Order Logic



R: Richard the Lionheart  
J: King John (his brother)



- Formally, a relation is the **set of tuples** of objects that are related
- Tuples are always **ordered**
- The “brother” **relation** in this model is the set  
 $\{ \langle \text{Richard the Lionheart}, \text{King John} \rangle, \langle \text{King John}, \text{Richard the Lionheart} \rangle \}$
- The “on head” **relation** in this model is the set  
 $\{ \langle \text{the crown}, \text{King John} \rangle \}$
- The “person” **property** (unary relation) is the set  
 $\{ \langle \text{Richard the Lionheart} \rangle, \langle \text{King John} \rangle \}$



- Functions relate **tuples** to **exactly one object**
- The (unary) “left leg” function includes:
  - $\langle \text{Richard the Lionheart} \rangle \rightarrow \text{Richard's left leg}$
  - $\langle \text{King John} \rangle \rightarrow \text{King John's left leg}$
- All functions in FOL must be **total functions**, i.e., there must be a value for every input tuple
- The crown also must have a left leg and so must each of the left legs.
- We introduce an “invisible” object that is the left leg of everything that has no left leg.
  - as long as we don't use it, there is no problem



- The **basic syntactic elements** of FOL are the symbols for objects, relations and functions
  - **constant symbols** for objects  
here: *Richard, John*
  - **predicate symbols** for relations  
here: *Brother, OnHead, Person, King, Crown*
  - **functions symbols** for functions  
here: *LeftLeg*
- Convention: Symbols start with **uppercase** letters




- Each model must have an **interpretation** that specifies exactly which objects, relations and functions are referred to by the symbols
- One possible (**intended**) interpretation would be
  - *Richard* refers to Richard the Lionheart
  - *John* refers to the evil King John
  - ...
- Another possible interpretation could be
  - *Richard* refers to the crown
  - *John* refers to King John's left leg

## 8.2 Syntax and Semantics of First-Order Logic



- Note that there are objects without a name, e.g., the intended interpretation does not name the crown or the legs.
- It is also possible that one object has two names



- Some members of the set of all models with two constant symbols,  $R$  and  $J$ . The interpretation of each constant symbol is shown by a gray arrow.



- A **term** is a **logical expression** that refers to an object, e.g., *LeftLeg(John)*
- “King John’s left leg” does not name the leg, but refers to it
- Note that this is **not a subroutine** call that returns a value, it is just like a complicated name
- We can **reason** about left legs, without ever providing a definition of *LeftLeg*



- Formal semantics:  
Consider a term  $f(t_1, \dots, t_n)$ 
  - The function symbol  $f$  refers to some function in the model (call it  $F$ )
  - The argument terms refer to objects in the domain (call them  $d_1, \dots, d_n$ )
  - The term as a whole refers to the object that is the value of the function  $F$  applied to  $d_1, \dots, d_n$ .
- If *LeftLeg* refers to the “left leg” function and *John* refers to King John, *LeftLeg(John)* refers to King John’s left leg.

## 8.2 Syntax and Semantics of First-Order Logic



- An **atomic sentence** (or **atom**) is formed from a **predicate symbol** optionally followed by a list of terms, such as
  - $Brother(Richard, John)$
  - $Married(Father(Richard), Mother(John))$
- We can use **logical connectives** to construct more complex sentences
  - $\neg Brother(LeftLeg(Richard), John)$
  - $Brother(Richard, John) \wedge Brother(John, Richard)$
  - $King(Richard) \vee King(John)$
  - $\neg King(Richard) \Rightarrow King(John)$

## 8.2 Syntax and Semantics of First-Order Logic

- **Quantifiers** let us express properties of entire collections of objects, instead of enumerating the objects by name
- **Universal** quantification:  $\forall x$
- **Existential** quantification:  $\exists x$
- $x$  is called a variable
- variables are terms by itself
- A term with no variables is called **ground term**



## 8.2 Syntax and Semantics of First-Order Logic

- For all  $x$ , if  $x$  is a king, then  $x$  is a person  
$$\forall x \text{ King}(x) \Rightarrow \text{Person}(x)$$
- If there are five objects, this universally quantified sentence is equivalent to asserting the following five sentences:
  - Richard is a king  $\Rightarrow$  Richard is a person
  - John is a king  $\Rightarrow$  John is a person
  - Richard's left leg is a king  $\Rightarrow$  Richard's left leg is a person
  - John's left leg is a king  $\Rightarrow$  John's left leg is a person
  - The Crown is a king  $\Rightarrow$  the crown is a person

## 8.2 Syntax and Semantics of First-Order Logic

- A **common mistake** is to use conjunction instead of implication

$$\forall x \text{ King}(x) \wedge \text{Person}(x)$$

- This sentence would also state that Richard's left leg is a king and Richard's left leg is a person, which is not what we wanted to express.

## 8.2 Syntax and Semantics of First-Order Logic

- There exists an  $x$  such that  $x$  is a crown and  $x$  is on the head of King John  
$$\exists x \text{ Crown}(x) \wedge \text{OnHead}(x, \text{John})$$
- Similar to universal quantification, this expressed that of all the objects in the model, at **least one** makes this assertion **true**.
- A **common mistake** here is to use the implication instead of the conjunction:  
$$\exists x \text{ Crown}(x) \Rightarrow \text{OnHead}(x, \text{John})$$
- If  $x$  is not a crown, the assertion is true, which is not what we wanted to express.

## 8.2 Syntax and Semantics of First-Order Logic

- Quantifiers can be **nested**

$$\forall x \forall y \text{ Brother}(x, y) \Rightarrow \text{Sibling}(x, y)$$

- But the **order** is important:

$$\forall x \exists y \text{ Loves}(x, y)$$

(Everybody loves somebody)

is not the same as

$$\exists y \forall x \text{ Loves}(x, y)$$

(There is someone who is loved by everyone)

## 8.2 Syntax and Semantics of First-Order Logic

- Because  $\forall$  is a conjunction over the universe of objects and  $\exists$  is a disjunction, they obey **De Morgan's rules**:

$$\forall x \neg P \equiv \neg \exists x P$$

$$\neg \forall x P \equiv \exists x \neg P$$

$$\forall x P \equiv \neg \exists x \neg P$$

$$\exists x P \equiv \neg \forall x \neg P$$

## 8.2 Syntax and Semantics of First-Order Logic



- We can use the **equality symbol** to signify that two terms refer to the same object

$$\textit{Father}(\textit{John}) = \textit{Henry}$$

- In the same way, we can express that two symbols refer to different objects

$$\begin{aligned} \exists x \exists y \textit{ Brother}(x, \textit{Richard}) \wedge \textit{ Brother}(y, \textit{Richard}) \\ \wedge \neg(x = y) \end{aligned}$$

- The above sentence states that Richard has at least two brothers.
- Sometimes,  $x \neq y$  is used for  $\neg(x = y)$

## 8.2 Syntax and Semantics of First-Order Logic

- There are **alternative semantics**, such as the **database semantics**, which has:
  - the **unique-names assumption**:  
We assume that different symbols refer to different objects.
  - the **closed-world assumption**:  
We assume that all atomic sentences not known to be true are in fact false.
  - **domain closure**:  
We assume that each model contains no more domain elements than those named by the constant symbols

- TELL adds **knowledge** to the **database**:

*TELL(KB, King(John))*

*TELL(KB, Person(Richard))*

*TELL(KB,  $\forall x \text{ King}(x) \Rightarrow \text{Person}(x)$ )*

- ASK **queries** it (exactly like in PL):

*ASK(KB, King(John))*                      returns true

*ASK(KB,  $\exists x \text{ Person}(x)$ )*                      returns true

- The last query is like asking “Can you tell me the time?” and getting the answer “Yes.”

*ASK(KB,  $\exists x \text{ Person}(x)$ )*                      returns  $\{x/\text{John}\}$   
and  $\{x/\text{Richard}\}$



- The kinship domain

- $\forall m, c \text{ Mother}(c) = m \Leftrightarrow \text{Female}(m) \wedge \text{Parent}(m, c)$
- $\forall w, h \text{ Husband}(h, w) \Leftrightarrow \text{Male}(h) \wedge \text{Spouse}(h, w)$
- $\forall x \text{ Male}(x) \Leftrightarrow \neg \text{Female}(x)$
- $\forall p, c \text{ Parent}(p, c) \Leftrightarrow \text{Child}(c, p)$
- $\forall g, c \text{ Grandparent}(g, c) \Leftrightarrow \exists p \text{ Parent}(g, p) \wedge \text{Parent}(p, c)$
- $\forall x, y \text{ Sibling}(x, y) \Leftrightarrow x \neq y \wedge \exists p \text{ Parent}(p, x) \wedge \text{Parent}(p, y)$
- ...

- Each of these sentences is an **axiom**, as they provide **basic factual information**
- They are **definitions**, since they have the form:
  - $\forall x, y \ P(x, y) \Leftrightarrow \dots$
- Some logical sentences are **theorems**
  - $\forall x, y \ Sibling(x, y) \Leftrightarrow Sibling(y, x)$
- Theorems provide **no new information**, as they are **entailed** by the **axioms**
- But they reduce the computational cost of deriving new sentences

## 8.3 Using First-Order Logic

---

- Not all axioms are definitions
- For example, we do not know enough to fully define what a person is
  - $\forall x \text{ Person}(x) \Leftrightarrow \dots$
- But we can still have axioms that are **partial specifications of properties**
  - $\forall x \text{ Person}(x) \Rightarrow \dots$
  - $\forall x \dots \Rightarrow \text{Person}(x)$
- Axioms can also be **plain facts**:
  - $\text{Male}(\text{Jim})$
  - $\text{Spouse}(\text{Jim}, \text{Laura})$

- The **Peano axioms** define natural numbers
  - $NatNum(0)$
  - $\forall n \text{ } NatNum(n) \Rightarrow NatNum(S(n))$
- With this definition, successors ( $S(n)$ ) of natural numbers are also natural numbers
  - Natural numbers are:  $0, S(0), S(S(0)), S(S(S(0))), \dots$
- We need to constrain the successor function
  - $\forall n \text{ } 0 \neq S(n)$
  - $\forall m, n \text{ } m \neq n \Rightarrow S(m) \neq S(n)$

- Now we can define addition
  - $\forall m \text{ NatNum}(m) \Rightarrow +(0, m) = m$
  - $\forall m, n \text{ NatNum}(m) \wedge \text{NatNum}(n) \Rightarrow +(S(m), n) = S(+(m, n))$
- For better readability, we can use the **infix** notation  $(m + n)$  instead of **prefix**  $(+(m, n))$ 
  - $\forall m, n \text{ NatNum}(m) \wedge \text{NatNum}(n) \Rightarrow (m+1)+n = (m+n)+1$
- The infix notation is **syntactic sugar**, that is, an extension to or abbreviation of the standard syntax, that does **not change semantics**.

- **Set theory**
  - The empty set is a constant written as  $\{ \}$
  - Binary predicates are  $x \in s$  and  $s_1 \subseteq s_2$
  - Binary functions are  $s_1 \cap s_2, s_1 \cup s_2, \{x/s\}$  (the set resulting from adjoining element  $x$  to set  $s$ )
- **Lists** are similar to sets, but are ordered
  - The empty list is a constant written as  $[]$  or *Nil*
  - The function  $Cons(x, Nil)$  is written as  $[x]$
  - and  $Cons(x, y)$ , with  $y$  being a list, as  $[x/y]$
  - A list of several elements like  $[A, B, C]$  corresponds to  $Cons(A, Cons(B, Cons(C, Nil)))$

The wumpus world:

- the agent receives a **percept vector** with five elements
- we must also include the **time** at which a percept occurred, for which we use integers
  - *Percept([Stench, Breeze, Glitter, None, None], 5)*
- The **actions** can be represented as **logical terms**
  - *Turn(Right), Turn(Left), Forward, Shoot, Grab, Climb*
- To determine which is best, the agents asks
  - *AskVars( $\exists a$  BestAction( $a, 5$ ))*
- This returns a **binding list** such as  $\{a/Grab\}$

- The raw percept data **implies** certain facts about the current state, for example:
  - $\forall t, s, g, m, c \text{ Percept}([s, \text{Breeze}, g, m, c], t) \Rightarrow \text{Breeze}(t)$
  - $\forall t, s, b, m, c \text{ Percept}([s, b, \text{Glitter}, m, c], t) \Rightarrow \text{Glitter}(t)$
- Simple **reflex** behavior can be represented as
  - $\forall t \text{ Glitter}(t) \Rightarrow \text{BestAction}(\text{Grab}, t)$
- Instead of naming squares, we identify them by their row and column
  - $\forall x, y, a, b \text{ Adjacent}([x, y], [a, b]) \Leftrightarrow$   
 $(x = a \wedge (y = b - 1 \vee y = b + 1)) \vee$   
 $(y = b \wedge (x = a - 1 \vee x = a + 1))$



- The agent's location changes over time, so
  - $At(Agent, s, t)$  means the agent is at square  $s$  at time  $t$
- The wumpus can't move, so we fix its position
  - $\forall t At(Wumpus, [2, 2], t)$
- Objects can only be at one location at a time
  - $\forall x, s_1, s_2, t At(x, s_1, t) \wedge At(x, s_2, t) \Rightarrow s_1 = s_2$
- Possible **inference rules** are
  - $\forall s, t At(Agent, s, t) \wedge Breeze(t) \Rightarrow Breezy(s)$
  - $\forall s Breezy(s) \Leftrightarrow \exists r Adjacent(r, s) \wedge Pit(r)$
  - $\forall t HaveArrow(t+1) \Leftrightarrow (HaveArrow(t) \wedge \neg Action(Shoot, t))$

## 8.4 Knowledge Engineering in First-Order Logic



- The process of constructing a knowledge-base is called **knowledge engineering**
- It can vary widely in content, scope, and difficulty, but it includes the following steps:
  1. Identify the task
  2. Assemble the relevant knowledge
  3. Decide on a vocabulary
  4. Encode general knowledge about the domain
  5. Encode a description of the specific problem instance
  6. Pose queries to the inference procedure
  7. Debug the knowledge base



- First-order logic is far **more powerful** than propositional logic
- Knowledge representation should be **declarative, compositional, expressive, context independent** and **unambiguous**
- Logics differ in their **ontological** and **epistemological commitments**
- The **syntax** of FOL builds on that of PL
- A **possible world**, or **model**, for FOL includes a set of objects and an **interpretation** that maps symbols to objects, predicates and functions